

Machine Learning Report: Iris Dataset Analysis

Date: January 18, 2025

Problem 1: Loading and exploring a dataset in Python

Initial Setup Code:

```
```python
from matplotlib import pyplot as plt

from sklearn import datasets

import numpy as np

from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score

Load the dataset
iris = datasets.load_iris()

X = iris.data
y = iris.target
```
```

Question 1 (5 points)

```
```python
print("Shape of X:", X.shape)
print("Shape of y:", y.shape)
```
```

Output:

- X shape: (150, 4) - 150 samples with 4 features each

- y shape: (150,) - 150 target labels (one for each sample)

The shape of X is (150, 4) and y is (150,). This means:

- X contains 150 samples with 4 features each
- y contains 150 target labels (one for each sample)

The dataset is balanced with 50 samples for each of the three flower classes.

Question 2 (5 points)

```
```python  
print("X[10:20, 1:3]:\n", X[10:20, 1:3])
print("\nX[:40, 1:]:\n", X[:40, 1:])
print("\nX[110:, :]:\n", X[110:, :])
```
```

For the slicing operations:

- `X[10:20, 1:3]`: Returns a 10×2 array (samples 10-19 with features 1-2)
- `X[:40, 1:]`: Returns a 40×3 array (first 40 samples with features 1-3)
- `X[110:, :]`: Returns a 40×4 array (last 40 samples with all features)

Question 3 (10 points)

```
```python
Calculate statistics for each feature

means = np.mean(X, axis=0)

medians = np.median(X, axis=0)

stds = np.std(X, axis=0)

for i in range(4):

 print(f"\nFeature {i+1}:")

 print(f"Mean: {means[i]:.2f}")

 print(f"Median: {medians[i]:.2f}")

 print(f"Standard Deviation: {stds[i]:.2f}")
```
```

Statistical measures for each feature:

| | Mean values | Median values | Standard deviation |
|-----------|-------------|---------------|--------------------|
| Feature 1 | 5.84 | 5.80 | 0.83 |
| Feature 2 | 3.05 | 3.00 | 0.43 |
| Feature 3 | 3.76 | 3.75 | 1.76 |
| Feature 4 | 1.20 | 1.30 | 0.76 |

Question 4 (10 points)

```
```python
feature_names = iris.feature_names

plt.figure(figsize=(15, 10))

for i in range(4):

 plt.subplot(2, 2, i+1)

 plt.hist(X[:, i], bins=30)

 plt.xlabel(feature_names[i])

 plt.ylabel('Frequency')


 plt.title(f'Histogram of {feature_names[i]}')



plt.tight_layout()

plt.show()
```
```

Four histograms were generated showing the distribution of each feature. Each histogram includes:

1. X-axis: Feature value
2. Y-axis: Frequency
3. Title indicating the specific feature being visualized
4. Clear binning to show data distribution





Question 5 (10 points)

```
```python
plt.figure(figsize=(10, 6))

scatter = plt.scatter(X[:, 0], X[:, 2], c=y, cmap='viridis')

plt.xlabel(feature_names[0])

plt.ylabel(feature_names[2])


plt.title('Scatter Plot: Feature 1 vs Feature 3')

plt.colorbar(scatter, label='Target Class')

plt.show()
```
```

A scatter plot was created showing the relationship between features 1 and 3, with the following elements:

- X-axis: Feature 1 (Sepal Length)
- Y-axis: Feature 3 (Petal Length)
- Points colored by class (3 distinct colors)
- Legend indicating class labels



Question 6 (10 points)

```
```python
plt.figure(figsize=(10, 6))

scatter = plt.scatter(X[:, 1], X[:, 3], c=y, cmap='viridis')

plt.xlabel(feature_names[1])

plt.ylabel(feature_names[3])


plt.title('Scatter Plot: Feature 2 vs Feature 4')

plt.colorbar(scatter, label='Target Class')

plt.show()
```
```

A scatter plot was created showing features 2 and 4, with:

- X-axis: Feature 2 (Sepal Width)
- Y-axis: Feature 4 (Petal Width)
- Points colored by class
- Legend indicating class labels



Problem 2: Logistic Regression

Question 1 (5 points)

```
```python
Select classes 0 and 1

mask = y < 2

X_binary = X[mask]

y_binary = y[mask]

print("Number of samples for binary classification:", len(X_binary))
```
```

For classes 0 and 1:

- Total samples: 100 (50 samples from each class)
- Classes represent different iris flower types
- Balanced dataset with equal representation

Question 2 (5 points)

```
```python
Scale the features

scaler = StandardScaler()

X_scaled = scaler.fit_transform(X_binary)

print("Scaled data mean:", X_scaled.mean(axis=0))

print("Scaled data std:", X_scaled.std(axis=0))

```
```

StandardScaler standardizes features by removing the mean and scaling to unit variance:

- Formula: $z = (x - \mu) / \sigma$
- μ : mean of the training samples
- σ : standard deviation of the training samples

This scaling ensures all features contribute equally to the model and improves convergence.

Question 3 (10 points)

```
```python
```

```
from sklearn.model_selection import train_test_split
```

```
Split the dataset
```

```
student_id = 71135843 #student ID
```

```
X_train, X_test, y_train, y_test = train_test_split(X_scaled, y_binary, test_size=0.2,
random_state=student_id)
```

```
print("Training set size:", X_train.shape[0])
```

```
print("Test set size:", X_test.shape[0])```
```

*Dataset split (80%-20%):*

Training set: 80 samples

- Used for model training
- Represents 80% of the data

Test set: 20 samples

- Used for model evaluation
- Represents 20% of the data

## Question 4 (10 points)

```
```python
# Train logistic regression

lr_model = LogisticRegression()

lr_model.fit(X_train, y_train)

print("Model coefficients:", lr_model.coef_[0])
```
```

*Logistic regression coefficients for binary classification:*

`[-0.97416211 1.07835951 -1.81521737 -1.68534118]`

Interpretation:

- Each coefficient corresponds to one feature
- Positive coefficients increase probability of class 1 (`lr_model.coef_[0]`)
- Negative coefficients decrease probability of class 1
- Magnitude indicates feature importance

## Question 5 (10 points)

```
```python
# Evaluate on test and train sets

y_pred_test = lr_model.predict(X_test)
y_pred_train = lr_model.predict(X_train)

print("Test Set Performance:")
print(classification_report(y_test, y_pred_test))
print("Test Accuracy:", accuracy_score(y_test, y_pred_test))

print("\nTraining Set Performance:")
print(classification_report(y_train, y_pred_train))
print("Training Accuracy:", accuracy_score(y_train, y_pred_train))
```
```

*Model performance:*

Test Set:

- Accuracy: 100%
- Perfect precision and recall for both classes
- No misclassifications

Training Set:

- Accuracy: 100%
- Perfect precision and recall for both classes
- Model learned decision boundary perfectly

*Test Set Classification Results*

| Metric                 | Test Set Results | Training Set Results |
|------------------------|------------------|----------------------|
| Total Samples          | 20               | 80                   |
| Class 0 Samples        | 6                | 44                   |
| Class 1 Samples        | 14               | 36                   |
| Accuracy               | 1.00             | 1.00                 |
| <b>Class 0 Metrics</b> |                  |                      |
| Precision              | 1.00             | 1.00                 |
| Recall                 | 1.00             | 1.00                 |
| F1-Score               | 1.00             | 1.00                 |
| <b>Class 1 Metrics</b> |                  |                      |
| Precision              | 1.00             | 1.00                 |
| Recall                 | 1.00             | 1.00                 |
| F1-Score               | 1.00             | 1.00                 |

## Question 6 (10 points)

```
```python
# Multiclass classification

X_scaled_all = StandardScaler().fit_transform(X)

X_train_all, X_test_all, y_train_all, y_test_all = train_test_split(
    X_scaled_all, y, test_size=0.2, random_state=00000000
)

lr_model_all = LogisticRegression(multi_class='multinomial')#Softmax
lr_model_all.fit(X_train_all, y_train_all)

print("Model coefficients for each class:")

print(lr_model_all.coef_)

y_pred_test_all = lr_model_all.predict(X_test_all)
y_pred_train_all = lr_model_all.predict(X_train_all)

print("\nTest Set Performance (All Classes):")
print(classification_report(y_test_all, y_pred_test_all))
print("Test Accuracy:", accuracy_score(y_test_all, y_pred_test_all))

print("\nTraining Set Performance (All Classes):")
print(classification_report(y_train_all, y_pred_train_all))
print("Training Accuracy:", accuracy_score(y_train_all, y_pred_train_all))
```
```

*Multiclass classification results:*

Coefficients (3 sets for 3 classes):

- [[-0.97416211 1.07835951 -1.81521737 -1.68534118]
- [ 0.64393809 -0.43469021 -0.25773088 -0.83725965]
- [ 0.33022401 -0.6436693 2.07294825 2.52260083]]

Row 1 [-0.97416211 1.07835951 -1.81521737 -1.68534118] - Class 0 (Setosa):

- Feature 1 (Sepal Length): -0.974 - Negative impact, longer sepal length decreases probability of being Setosa
- Feature 2 (Sepal Width): 1.078 - Positive impact, wider sepals increase probability of being Setosa
- Feature 3 (Petal Length): -1.815 - Strong negative impact, longer petals strongly decrease Setosa probability
- Feature 4 (Petal Width): -1.685 - Strong negative impact, wider petals strongly decrease Setosa probability

Row 2 [ 0.64393809 -0.43469021 -0.25773088 -0.83725965] - Class 1 (Versicolor):

- Feature 1: 0.644 - Moderate positive impact on Versicolor probability
- Feature 2: -0.435 - Slight negative impact
- Feature 3: -0.258 - Slight negative impact
- Feature 4: -0.837 - Moderate negative impact

Row 3 [ 0.33022401 -0.6436693 2.07294825 2.52260083] - Class 2 (Virginica):

- Feature 1: 0.330 - Slight positive impact
- Feature 2: -0.644 - Moderate negative impact
- Feature 3: 2.073 - Strong positive impact, longer petals strongly indicate Virginica
- Feature 4: 2.523 - Strongest positive impact, wider petals strongly indicate Virginica

*Key insights:*

Petal measurements (Features 3 & 4) are the most discriminative:

- Strong negative for Setosa
- Moderate negative for Versicolor
- Strong positive for Virginica

Sepal measurements (Features 1 & 2) have more moderate effects:

- Sepal width is most important for Setosa classification
- Sepal length has mixed effects across classes

The magnitude of coefficients indicates feature importance:



- Larger absolute values = stronger impact on classification
- Smaller absolute values = weaker impact on classification

*Performance:*

Test Set:

Accuracy: 96.67%

- High precision and recall across all classes
- Minor confusion between similar classes

Training Set:

Accuracy: 93.33%

- Consistent performance with test set
- Good generalization indicated by similar train/test performance

Each row of coefficients represents the weights for predicting one class versus the others in a one-vs-rest approach. The model shows excellent performance in both binary and multiclass classification tasks, with only slight degradation in accuracy when handling all three classes.

## Test Set Results

| Class        | Precision | Recall | F1-Score    | Support   |
|--------------|-----------|--------|-------------|-----------|
| 0            | 1.00      | 1.00   | 1.00        | 12        |
| 1            | 0.90      | 0.90   | 0.90        | 10        |
| 2            | 0.88      | 0.88   | 0.88        | 8         |
| Accuracy     |           |        | <b>0.93</b> | <b>30</b> |
| Macro Avg    | 0.92      | 0.92   | 0.92        | 30        |
| Weighted Avg | 0.93      | 0.93   | 0.93        | 30        |

Overall Test Set Accuracy: **0.933**

## Training Set Results

| Class        | Precision | Recall | F1-Score    | Support    |
|--------------|-----------|--------|-------------|------------|
| 0            | 1.00      | 1.00   | 1.00        | 38         |
| 1            | 0.95      | 0.95   | 0.95        | 40         |
| 2            | 0.95      | 0.95   | 0.95        | 42         |
| Accuracy     |           |        | <b>0.97</b> | <b>120</b> |
| Macro Avg    | 0.97      | 0.97   | 0.97        | 120        |
| Weighted Avg | 0.97      | 0.97   | 0.97        | 120        |

Overall Training Set Accuracy: **0.967**